

# **SESSION 11**

# **DATA WRANGLING 1**

R FOR SOCIAL DATA SCIENCE

JEFFREY ZIEGLER, PHD

ASSISTANT PROFESSOR IN POLITICAL SCIENCE & DATA SCIENCE  
TRINITY COLLEGE DUBLIN

FALL 2022

# ROAD MAP FOR TODAY

Last time:

- Debugging
- Handling conditions
- Testing
- Defensive programming

This time:

- Data frames in base R
- Alternatives to data frames

# DATA ORGANIZATION

■ 'Tidy' data is a specific subset of rectangular data, where:

- ▶ Each variable is in a column
- ▶ Each observation is in a row
- ▶ Each value is in a cell

country	year	cases	population
Afghanistan	1999	18215	19995071
Afghanistan	2000	18666	20095360
Brazil	1999	30737	17206362
Brazil	2000	80488	174504898
China	1999	210258	1272015272
China	2000	210766	128000583

variables

country	year	cases	population
Afghanistan	1999	18215	19995071
Afghanistan	2000	18666	20095360
Brazil	1999	30737	17206362
Brazil	2000	80488	174504898
China	1999	210258	1272015272
China	2000	210766	128000583

observations

country	year	cases	population
Afghanistan	1999	18215	19995071
Afghanistan	2000	18666	20095360
Brazil	1999	30737	17206362
Brazil	2000	80488	174504898
China	1999	210258	1272015272
China	2000	210766	128000583

values

Source: R for Data Science

# DATA FRAMES

- Data frame is one of the object types available in base R
- Despite their matrix-like appearance, data frames are lists of equal-sized vectors
- Data frames can be created with 'data.frame()' function with named vectors as input

```
1 df <- data.frame(  
2   x = 1:4,  
3   y = c("a", "b", "c", "d"),  
4   z = c(TRUE, FALSE, FALSE, TRUE)  
5 )
```

	x	y	z
1	1	a	TRUE
2	2	b	FALSE
3	3	c	FALSE
4	4	d	TRUE

## EXAMPLE DATA FRAME

```
1 # str() function applied to data frame is useful in determining  
  variable types
```

```
2 str(df)
```

```
'data.frame': 4 obs. of 3 variables:
```

```
$ x: int  1 2 3 4
```

```
$ y: chr  "a" "b" "c" "d"
```

```
$ z: logi  TRUE FALSE FALSE TRUE
```

```
1 # dim() function behaves similar to matrix, showing N rows and N  
  columns, respectively
```

```
2 dim(df)
```

```
[1] 4 3
```

```
1 # In contrast to matrix length() of data frame displays the length  
  of underlying list
```

```
2 length(df)
```

```
[1] 4 3
```

# CREATING DATA FRAME

```
1 l <- list(x = 1:5, y = letters[1:5], z = rep(c(TRUE, FALSE),  
2     length.out = 5))  
l
```

```
$x  
[1] 1 2 3 4 5
```

```
$y  
[1] "a" "b" "c" "d" "e"
```

```
$z  
[1] TRUE FALSE TRUE FALSE TRUE
```

```
1 df <- data.frame(l)  
2 df
```

```
x y      z  
1 1 a  TRUE  
2 2 b FALSE  
3 3 c  TRUE  
4 4 d FALSE  
5 5 e  TRUE
```

## SUBSETTING DATA FRAME

- In subsetting data frames, techniques of subsetting matrices and lists are combined
- If you subset with a single vector, it behaves as a list
- If you subset with two vectors, it behaves as a matrix

## EX: SUBSETTING DATA FRAME

```
1 # Like a list
2 df[c("x", "z")]
```

```
  x      z
1 1  TRUE
2 2 FALSE
3 3  TRUE
4 4 FALSE
5 5  TRUE
```

```
1 # Like a matrix
2 df[, c("x", "z")]
```

```
  x      z
1 1  TRUE
2 2 FALSE
3 3  TRUE
4 4 FALSE
5 5  TRUE
```

```
1 df[df$y=="b",]
```

```
  x y      z
2 2 b FALSE
```



## BUILDING DATA FRAME

- `'rbind()'` (row bind) - appends a row to data frame
- `'cbind()'` (column bind) - appends a column to data frame
- Both require compatible sizes (number of rows/columns)

# BUILDING DATA FRAME: ADDING COLUMNS

```
1 set.seed(12345)
2 rand <- rnorm(5)
3 rand
```

```
[1] 0.5855288 0.7094660 -0.1093033 -0.4534972 0.6058875
```

```
1 df <- cbind(df, rand)
2 df
```

	x	y	z	rand
1	1	a	TRUE	0.5855288
2	2	b	FALSE	0.7094660
3	3	c	TRUE	-0.1093033
4	4	d	FALSE	-0.4534972
5	5	e	TRUE	0.6058875

# BUILDING DATA FRAME: ADDING ROWS

```
1 # Note, row has to be list since contains different data types
2 r <- list(6, letters[6], FALSE, rnorm(1))
3 r
```

```
[[1]]
[1] 6
```

```
[[2]]
[1] "f"
```

```
[[3]]
[1] FALSE
```

```
[[4]]
[1] -1.817956
```

```
1 df <- rbind(df, r)
2 df
```

```
x y      z      rand
1 1 a  TRUE  0.5855288
2 2 b FALSE  0.7094660
3 3 c  TRUE -0.1093033
4 4 d FALSE -0.4534972
5 5 e  TRUE  0.6058875
6 6 f FALSE -1.8179560
```

# ADDING/MODIFYING COLUMNS IN BASE R

```
1 # New columns can be created/modified by assignment (if RHS object  
  has correct length)
```

```
2 df["r"] <- rnorm(6)
```

```
3 df
```

	x	y	z	rand	r
1	1	a	TRUE	0.5855288	0.6300986
2	2	b	FALSE	0.7094660	-0.2761841
3	3	c	TRUE	-0.1093033	-0.2841597
4	4	d	FALSE	-0.4534972	-0.9193220
5	5	e	TRUE	0.6058875	-0.1162478
6	6	f	FALSE	-1.8179560	1.8173120

```
1 # Individual columns can also be selected with $ operator
```

```
2 df$r <- df$r + 5
```

```
3 df
```

	x	y	z	rand	r
1	1	a	TRUE	0.5855288	5.630099
2	2	b	FALSE	0.7094660	4.723816
3	3	c	TRUE	-0.1093033	4.715840
4	4	d	FALSE	-0.4534972	4.080678
5	5	e	TRUE	0.6058875	4.883752
6	6	f	FALSE	-1.8179560	6.817312

# RENAMING COLUMNS IN BASE R

```
1 # names() attribute for data frames/tibbles contains column names
2 names(df)
```

```
[1] "x"      "y"      "z"      "rand" "r"
```

```
1 # Individual columns can also be selected with $ operator
2 df$r <- df$r + 5
3 df
```

	x	y	z	rand_new	r
1	1	a	TRUE	0.5855288	5.630099
2	2	b	FALSE	0.7094660	4.723816
3	3	c	TRUE	-0.1093033	4.715840
4	4	d	FALSE	-0.4534972	4.080678
5	5	e	TRUE	0.6058875	4.883752
6	6	f	FALSE	-1.8179560	6.817312

# ISSUES WITH DATA FRAME

- While very versatile (and available out-of-the-box) data frames have their drawbacks:
  - ▶ Individual cells (observations) cannot themselves be lists
  - ▶ Somewhat limited (and inconsistent) data manipulation functions
  - ▶ Memory inefficient (*copy-on-modify* semantics)
  - ▶ No parallelisation

What's helpful for this? We'll talk about alternatives next time!

# DATA FORMATS IN R

- `'csv'` (Comma-separated value) files for storing tabular data
  - ▶ Standard file format for storing data that is highly interoperable across systems
  - ▶ Human-readable and can be opened in a simple text processor
- `'rds'` (R data serialization) files allow to store single R object
  - ▶ Can store arbitrary R objects (e.g. fitted model), similar to Python `'pickle'`
  - ▶ Offers data compression
  - ▶ Only works within R

# DATA FORMATS IN R

- 'rda' (R data) files for saving and loading multiple R objects
  - ▶ Offers data compression
  - ▶ Compares unfavourably to rds and, generally, should be avoided
- 'feather'/'parquet' - big data formats associated with **Apache Hadoop** ecosystem
  - ▶ Cutting-edge performance (compression and read/write access)
  - ▶ Not human-readable
  - ▶ Relatively new, could be an overkill for some tasks



# FUNCTIONS FOR DATA I/O

## ■ `'csv'` (Comma-separated value)

- ▶ `'read.csv()'`/`'write.csv()'` - base R functions
- ▶ `'readr::read_csv()'`/`'readr::write_csv()'` - functions from `'readr'` package in `'tidyverse'`

## ■ `'rds'` (R data serialization)

- ▶ `'readRDS()'`/`'writeRDS()'` - base R functions
- ▶ `'readr::read_rds()'`/`'readr::write_rds()'` - functions from `'readr'` (no default compression)

# FUNCTIONS FOR DATA I/O

- `'rda'` (R data)

- ▶ `'save()'`/`'load()'` - base R functions

- `'feather'`/`'parquet'`

- ▶ `'arrow::read_feather()'`/`'arrow::write_feather()'` - functions from
  - ▶ `'arrow::read_parquet()'`/`'arrow::write_parquet()'` - `'arrow'` package in **Apache Arrow**

# READING DATA IN R EXAMPLE

```
1 # assuming your local GitHub is up-to-date
2 inc_local <- read.csv("../datasets/incumbents_subset.csv")
3 inc_url <- read.csv("https://raw.githubusercontent.com/
  jeffreyziegler/R_social_DS/main/datasets/incumbents_subset.csv
  ")
4 head(inc_local); head(inc_url)
```

	X	x	year	congress	chalspend	incspend	difflog	presvote	voteshare	inparty	incparty	seniority
1	53	53	1978	95	11.67655	12.24663	0.5700871	0.5267782	0.6023614	1	1	
2	54	54	1978	95	11.62039	12.49136	0.8709687	0.5659233	0.5836368	0	0	
3	55	55	1978	95	12.30557	12.73226	0.4266895	0.4646196	0.5922578	1	1	
4	56	56	1978	95	10.54843	12.50500	1.9565633	0.5012287	0.6992224	1	1	
5	57	57	1978	95	12.10366	12.76171	0.6580556	0.4774266	0.6419783	1	1	
6	58	58	1980	96	12.48744	12.83441	0.3469714	0.5901939	0.6257710	0	0	

# TUTORIAL - CREATING & MANIPULATING DATA FRAMES

- Load 'kaggle\_survey\_2021\_responses.csv' dataframe from GitHub repository to global environment
- Subset dataframe to include 'time to complete' and first 5 Qs
- Subset to women who have earned less than a doctoral degree

## TUTORIAL - DUMMY VARIABLES

- When analysing categorical data (particularly using it as independent variables in regression) it is common to construct dummy variables
  - Where categorical variables are represented by 1's and 0's depending on whether it is true or not for a given observation
  - For example, gender of respondents can be represented by 1's that indicate whether a given respondent is female (baseline/reference category) and 0's if they are not
1. Make a dummy variable for your two criteria from above (women and less than doctoral degrees)
  2. Subset original dataset based on new dummy variables

# OVERVIEW

This time:

- Data frames in base R
- Data input and output

Next time:

- Alternatives to data frames
- 'tidyverse' packages
- Working with tabular data
- Summary statistics